

Computer Organization and Operating System


Input and Output

Akharin Khunkitti

KMITL



Topic

- Input / Output Overview
- I/O Techniques
 - Programmed, Interrupt, DMA
- DMA – Direct Memory Access
- I/O Interfacing
- Conclusion




Input / Output Overview

- I/O Functions
 - Move/Transfer Data between computer and “External Devices”
 - Input - Move/Transfer Data From Outside to CPU/Memory
 - Output – Move/Transfer Data From CPU/Memory to Outside
- Internal Locations
 - CPU – Direct to Registers
 - Small number of Registers
 - Memory – Temporary places, wait for CPU processing
 - Can be big amount of data
- Normally Transfer between Main Memory
 - Input => Memory
 - Memory => Output




External Devices

- Human readable
 - Screen, printer, keyboard
- Machine readable
 - Monitoring and control
- Communication
 - Modem
 - Network Interface Card (NIC)




Input / Output Modules

- Wide variety of peripherals
 - Delivering different amounts of data
 - At different speeds
 - In different formats
- All slower than CPU and Memory (RAM)
- Need I/O modules
- Interface to CPU and Memory
- Interface to one or more peripherals




I/O Module Functions

- Control & Timing
- CPU Communication
- Device Communication
- Data Buffering
- Error Detection




I/O Transfer Steps

- CPU checks I/O module device status
- I/O module returns status
- If ready, CPU requests data transfer
- I/O module gets data from device
- I/O module transfers data to CPU
- Variations for output, DMA, etc.




Input / Output Techniques

- Programmed I/O
- Interrupt I/O
- Direct Memory Access (DMA)



Programmed I/O

- CPU has direct control over I/O
 - Sensing status
 - Read/write commands
 - Transferring data
- CPU waits for I/O module to complete operation
- Wastes CPU time
- Steps
 - CPU requests I/O operation
 - I/O module performs operation
 - I/O module sets status bits
 - CPU checks status bits periodically
 - I/O module does not inform CPU directly
 - I/O module does not interrupt CPU
 - CPU may wait or come back later



I/O Operations


- I/O Commands
 - CPU issues address
 - Identifies module (& device if >1 per module)
 - CPU issues command
 - Control - telling module what to do
 - e.g. spin up disk
 - Test - check status
 - e.g. power? Error?
 - Read/Write
 - Module transfers data via buffer from/to device
- Addressing I/O Devices
 - Under programmed I/O data transfer is very like memory access (CPU viewpoint)
 - Each device given unique identifier
 - CPU commands contain identifier (address)

MODES OF TRANSFER - PROGRAM-CONTROLLED I/O -

3 different Data Transfer Modes between the central computer(CPU or Memory) and peripherals;

Program-Controlled I/O
Interrupt-Initiated I/O
Direct Memory Access (DMA)

Program-Controlled I/O(Input Dev to CPU)



Polling or Status Checking

- Continuous CPU involvement
- CPU slowed down to I/O speed
- Simple
- Least hardware


I/O Mapping

- Memory mapped I/O
 - Devices and memory share an address space
 - I/O looks just like memory read/write
 - No special commands for I/O
 - Large selection of memory access commands available

- Isolated I/O
 - Separate address spaces
 - Need I/O or memory select lines
 - Special commands for I/O
 - Limited set

ADDRESS	INSTRUCTION	OPERAND	COMMENT
200	Load AC	"1"	Load accumulator
	Store AC	517	Initiate keyboard read
202	Load AC	517	Get status byte
	Branch if Sign = 0	202	Loop until ready
	Load AC	516	Load data byte

(a) Memory-mapped I/O





(b) Isolated I/O

ADDRESS	INSTRUCTION	OPERAND	COMMENT
200	Load I/O	5	Initiate keyboard read
201	Test I/O	5	Check for completion
	Branch Not Ready	201	Loop until complete
	In	5	Load data byte


Interrupt Driven I/O

- Overcomes CPU waiting
 - No repeated CPU checking of device
 - I/O module interrupts when ready
 - Basic Steps
 - CPU issues read command
 - I/O module gets data from peripheral
 - while CPU does other work
 - I/O module interrupts CPU
 - CPU requests data
 - I/O module transfers data




CPU Viewpoint for Interrupt

- Issue read command
- Do other work
- Check for interrupt at end of each instruction cycle
- If interrupted:-
 - Save context (registers)
 - Process interrupt
 - Fetch data & store





(a) Interrupt occurs after instruction
at location N



(b) Return from interrupt



Interrupt I/O

- Interrupt Module
- Different line for each module
 - PC
 - Limits number of devices
- Software poll
 - CPU asks each module in turn
 - Slow
- Multiple interrupts
 - i.e. an interrupt handler being interrupted
 - Each interrupt line has a priority
 - Higher priority lines can interrupt lower priority lines



DMA - Direct Memory Access

- Interrupt driven and programmed I/O require active CPU intervention
 - Transfer rate is limited
 - CPU is tied up
- DMA is the answer
- DMA Functions
 - Additional Module (hardware) on bus
 - DMA Controller
 - DMA controller takes over from CPU for I/O



DMA Operation

- CPU tells DMA controller:-
 - Read/Write
 - Device address
 - Starting address of memory block for data
 - Amount of data to be transferred
- CPU starts DMA Controller
 - CPU carries on with other work
- DMA controller deals with transfer
- DMA controller sends interrupt when finished









Figure 5-4. Operation of a DMA transfer.




DMA Transfer Cycle – Bus Stealing

- DMA controller takes over bus for a cycle
 - DMA Bus request to CPU
 - CPU sends DMA Bus Acknowledge to DMA Controller
 - When CPU not uses bus
 - DMA Controller can control the bus
 - Transfer of one word of data
- Not an interrupt
 - CPU does not switch context
- CPU suspended just before it accesses bus
 - i.e. before an operand or data fetch or a data write
- Slows down CPU but not as much as CPU doing transfer




DMA Fly-By

- While DMA using buses processor idle
- Processor using bus, DMA idle
 - Known as fly-by DMA controller
- Data does not pass through and is not stored in DMA chip
 - DMA only between I/O port and memory
 - Not between two I/O ports or two memory locations
- Can do memory to memory via register
- DMA Controller may contain multiple DMA channels
 - Programmed independently
 - Any one active
 - Numbered 0, 1, 2, 3, and ...





DMA Aside

- What effect does caching memory have on DMA?
 - Cache is a copy of memory
- What about on board cache?
- Hint:
 - How much are the system buses available?
 - cache controller support DMA




DMA Configurations

- Single Bus, Detached DMA controller
 - Each transfer uses bus twice
 - I/O to DMA then DMA to memory
 - CPU is suspended twice
- Single Bus, Integrated DMA controller
 - Controller may support >1 device
 - Each transfer uses bus once
 - DMA to memory
 - CPU is suspended once
- Separate I/O Bus
 - Bus supports all DMA enabled devices
 - Each transfer uses bus once
 - DMA to memory
 - CPU is suspended once




(b) Single-bus, Integrated DMA-I/O




(c) I/O bus

I/O Channels

- I/O devices getting more sophisticated
 - e.g. 3D graphics cards
- CPU instructs I/O controller to do transfer
- I/O controller does entire transfer
- Improves speed
 - Takes load off CPU
 - Dedicated processor is faster







(a) Selector



(b) Multiplexor

I/O Interfacing






- Connecting devices together
- Bit of wire?
- Dedicated processor/memory/buses?
- Hardware with Software
- I/O Interface Classification
 - Parallel vs Serial
 - Functions
 - Disk
 - Display
 - Keyboard/Mouse
 - Universal
 - Etc.
 - Internal vs External
 - Distances vs Device Number



Example: IEEE 1394 FireWire



- High performance serial bus
- Fast
- Low cost
- Easy to implement
- Also being used in digital cameras, VCRs and TV
- Configuration
 - Daisy chain
 - Up to 63 devices on single port
 - Really 64 of which one is the interface itself
 - Up to 1022 buses can be connected with bridges
 - Automatic configuration
 - No bus terminators
 - May be tree structure

Simple FireWire Configuration




Example: InfiniBand

- I/O specification aimed at high end servers
 - Merger of Future I/O (Cisco, HP, Compaq, IBM) and Next Generation I/O (Intel)
- Version 1 released early 2001
- Architecture and spec. for data flow between processor and intelligent I/O devices
- Intended to replace PCI in servers
- Increased capacity, expandability, flexibility
- Architecture
 - Remote storage, networking and connection between servers
 - Attach servers, remote storage, network devices to central fabric of switches and links
 - Greater server density
 - Scalable data centre
 - Independent nodes added as required
 - I/O distance from server up to
 - 17m using copper
 - 300m multimode fibre optic
 - 10km single mode fibre
 - Up to 30Gbps



Conclusion

- Input / Output Transfer Data Into or out from Computer Systems => Memory
- **Input / Output Techniques**
 - Programmed I/O – CPU Intensive - Slow
 - Interrupt I/O – CPU Faster
 - DMA – Fastest with HW Controller
- I/O Interfaces connected I/O Modules to computer system
- Many I/O Interfacing Standards



END

Questions?